

Geoff Keeling

📍 London, United Kingdom
💡 Ethics of AI/ML, Moral & Political Philosophy,
Decision Theory, Bioethics, Cognitive Science

in [/in/geoff-keeling](#)
🔗 [geoffkeeling.github.io](#)
✉ geoffkeelingbhyat@gmail.com

Employment

Staff Research Scientist , <i>Google Research</i> .	2025–Present
Senior Research Scientist , <i>Google Research</i> .	2023–25
Research Scientist , <i>Google Research</i> .	2023–23
Bioethicist , <i>Google Health</i> .	2021–23
Postdoctoral Fellow , Institute for Human-Centered AI & Center for Ethics in Society, <i>Stanford University</i> .	2020–21
Research Assistant , Leverhulme Centre for the Future of Intelligence, <i>University of Cambridge</i> .	2019–20

Affiliations

Fellow , Institute of Philosophy, School of Advanced Study, <i>University of London</i> .	2024–Present
Associate Fellow , Leverhulme Centre for the Future of Intelligence, <i>University of Cambridge</i> .	2020–Present

Education

PhD Philosophy , <i>University of Bristol</i> , No Corrections.	2017–20
MA Philosophy , <i>University of Bristol</i> , Distinction in all Components (Rank: 1st/24).	2016–17
BSc Philosophy, Logic, and Scientific Method , <i>London School of Economics</i> , First Class (Rank: Joint 1st/51).	2013–16

Book

Keeling, G. & Street, W. Groundwork for the science of AI welfare. Cambridge University Press. (Under Contract)

Papers

1. Gabriel, I. & **Keeling, G.** (2025). A matter of principle? AI alignment as the fair treatment of claims. *Philosophical Studies*, 1-23.
2. **Keeling, G.** & Street, W. (2025). On the attribution of confidence to large language models. *Inquiry*, 1-27.
3. Broestl, N.*, Lange, B.*, Voinea, C.*, **Keeling, G.***, & Lam, R. (2025). Evaluating intra-firm LLM alignment strategies in business contexts. *arXiv preprint arXiv:2505.18779v1*. (*Equal Contribution) [Media: [AI Insider](#)]
4. **Keeling, G.***, Street, W.*, Stachaczyk, M., Zakharova, D., Comşa, I., Sakovych, A., Logothetis, I., Zhang, Z., Agüera y Arcas, B., & Birch, J. (2024). Can LLMs make trade-offs involving stipulated pain and pleasure states? *arXiv preprint arXiv:2411.02432v1*. (*Equal Contribution) [Media: [Scientific American](#), [Futurism](#), [The Guardian](#), [MSN](#)]
5. Gabriel, I.*, Manzini, A.*, **Keeling, G.***, ... , Agüera y Arcas, B., Isaac, W., & Manyika, J. (2024). The ethics of advanced AI assistants. *arXiv preprint arXiv:2404.16244*. (*Equal Contribution) [Media: [WIRED](#), [CBS](#), [The Information](#), [Axios](#), [The Verge](#), [MSN](#), [Tech Policy Press](#), [MarkTechPost](#)]

6. Manzini, A., **Keeling, G.**, Marchal, N., McKee, K., Reiser, V., Haas, J., & Gabriel, I. (2024, June). Should users trust advanced AI assistants? Justified trust as a function of competence and alignment. In: *Proceedings of FAccT '24: ACM Conference on Fairness, Accountability, and Transparency*.
7. Manzini, A., **Keeling, G.**, Alberts, L., Vallor, S., Ringel Morris, M., & Gabriel, I. (2024, October). The code that binds us: Navigating the appropriateness of human-AI assistant relationships. In: *Proceedings of the Seventh AAAI/ACM Conference on AI, Ethics, and Society*. [🏆 Runner Up: Best Paper Award]
8. El-Sayed, S., Akbulut, C., McCroskery, A., **Keeling, G.**, Kenton, Z., Jalan, Z., Marchal, N., Manzini, A., Shevlane, T., Vallor, S., Susser, D., Franklin, M., Bridgers, S., Law, H., Rahtz, M., Shanahan, M., Tessler, M.H., Douillard, A., Everitt, T., & Brown, S. (2024). A mechanism-based approach to mitigating harms from persuasive generative AI. *arXiv preprint arXiv:2404.15058*. [Media: [VentureBeat](#), [Futurism](#)]
9. Street, W., Siy, J.O., **Keeling, G.**, Baranes, A., Barnett, B., McKibben, M., Kanyere, T., Lentz, A., Agüera y Arcas, B., & Dunbar, R. (2024). LLMs achieve adult human performance on higher-order theory of mind tasks. *arXiv preprint arXiv:2405.18870v1*.
10. Schaekermann, M., Spitz, T., Pyles, M., Cole-Lewis, H., Wulczyn, E., Pfhol, S.R., Martin, D., Jaroensri, R., **Keeling, G.**, Liu, Y., Farquhar, S., Xue, Q., Lester, J., Hughes, C., Strachan, P., Tan, F., Bui, P., Mermel, C.H., Peng, L.H., Matias, Y., Corrado, G.S., Webster, D.R., Virmani, S., Semturs, C., Liu, Y., Horn, I., & Chen, P-H.C. (2024). Health equity assessment of machine learning performance (HEAL): a framework and dermatology AI model case study. *Lancet eClinicalMedicine*. [Media: [MarkTechPost](#)]
11. Alberts, L., **Keeling, G.**, & McCroskery, A. (2024). Should agentic conversational AI change how we think about ethics? Characterising an interactional ethics centred on respect. *arXiv preprint arXiv:2401.09082v2*.
12. **Keeling, G.**, Lange, B., McCroskery, A., Pedersen, K., Weinberger, D., & Zevenbergen, B. (2024). Moral imagination for engineering teams: The technomoral scenario. *International Review of Information Ethics*, 34(1).
13. Lange, B.*, **Keeling, G.***, McCroskery, A., Zevenbergen, B., Blascovich, S., Pedersen, K., Lentz, A., & Agüera y Arcas, B. (2023). Engaging engineering teams through moral imagination: a bottom-up approach for responsible innovation and ethical culture change in technology companies. *AI and Ethics*, 1-10. (*Equal Contribution)
14. **Keeling, G.** (2023). Algorithmic bias, generalist models, and clinical medicine. *AI and Ethics*, 1-12.
15. **Keeling, G.** & Nyrupe R. Explainable machine learning, patient autonomy, and clinical reasoning. In: Véliz, C. (Ed). (2023). *Oxford Handbook of Digital Ethics* (pp. 528–552). Oxford University Press.
16. **Keeling, G.** (2023). A dilemma for reasons additivity. *Economics & Philosophy*, 39(1), 20-42.
17. Grote, T. & **Keeling, G.** (2022). Enabling fairness in healthcare through machine learning. *Ethics and Information Technology*, 24(3), 39.
18. **Keeling, G.** Automated vehicles and the ethics of classification. In: Jenkins, R., Cerný, D., & Hřibek, T. (Eds). (2022). *Autonomous Vehicle Ethics: The Trolley Problem and Beyond* (pp. 41-57). Oxford University Press.
19. Grote, T. & **Keeling, G.** (2022). On algorithmic fairness in medical practice. *Cambridge Quarterly of Healthcare Ethics*, 31(1), 83-94.
20. **Keeling, G.** & Burr, C. Digital manipulation and mental integrity. In: Jongepier, F., & Klenk, M. (2022). *The Philosophy of Online Manipulation* (pp. 253-271). Routledge.
21. **Keeling, G.** & Paterson, N. (2022). Proper functions: etiology without typehood. *Biology & Philosophy*, 37(3), 19.
22. Bommasani, R., ... , **Keeling, G.**, ... , & Liang, P. (2021). On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*. [Media: [Forbes](#), [The Economist](#), [VentureBeat](#), [WIRED](#)]
23. **Keeling, G.** (2020). Why trolley problems matter for the ethics of automated vehicles. *Science and Engineering Ethics*, 26(1), 293-307.

24. **Keeling, G.**, Evans, K., Thornton, S. M., Mecacci, G., & Santoni de Sio, F. Four perspectives on what matters for the ethics of automated vehicles. In: Meyer, G. & Beiker, S. (Eds). (2019). *Road Vehicle Automation 6* (pp. 49-60). Springer International Publishing.
25. **Keeling, G.** (2018). Legal necessity, Pareto efficiency & justified killing in autonomous vehicle collisions. *Ethical Theory and Moral Practice*, 21(2), 413-427.
26. **Keeling, G.** Against Leben's Rawlsian collision algorithm for autonomous vehicles. In: Müller, V. (Ed). (2018). *Philosophy and Theory of Artificial Intelligence 2017* (pp. 259-272). Springer International Publishing.
27. **Keeling, G.** (2018). The sensitivity argument against child euthanasia. *Journal of Medical Ethics*, 44(2), 143-144.
28. **Keeling, G.** (2018). Autonomy, nudging and post-truth politics. *Journal of Medical Ethics*, 44(10), 721-722.
29. **Keeling, G.** (2017). Commentary: Using virtual reality to assess ethical decisions in road traffic scenarios: applicability of value-of-life-based models and influences of time pressure. *Frontiers in Behavioral Neuroscience*, 11, 247-48.

Talks

1. **Panel: AI and Moral Patienthood**, with Henry Shevlin and Winnie Street.
AI, Animals and Digital Minds Conference, London, 2025.
2. **Animal-inspired Approaches to Assessing LLM Sentience**, with Winnie Street.
AI Welfare Workshop, Eleos AI, Berkeley, 2025.
3. **Ethics and the Agentic Turn**.
AI Ethics, Risks and Safety Conference, Bristol, 2025.
4. **What About the Moral Claims of AI Agents?**.
Symposium: Themes from Iason Gabriel, Department of Philosophy, Hong Kong University, 2025.
5. **Do Large Language Models Have Credences?**
AI and Ethics Workshop, Princeton University, 2024.
Institute for Ethics in Technology, Hamburg University of Technology, 2024.
6. **Can LLMs Make Trade-Offs Involving Stipulated Pain and Pleasure States?** with Winnie Street.
Institute of Philosophy, School of Advanced Study, University of London, 2024.
7. **The Ethics of Advanced AI Assistants**.
Being Human in the Age of AI, Institute of Philosophy, 2024.
Conference on the Ethics of Generative AI & Conversational Agents, LMU Munich, 2024.
8. **Educational Impact of Advanced AI Assistants**.
AI in Education Workshop, University of Bristol, 2024.
9. **Autonomy, Advice-Giving, and AI Assistants**.
Autonomy in an AI World, Imperial College London, 2024.
10. **Engaging Google Teams Through Moral Imagination**, with Amanda McCroskery, Ben Zevenbergen, et al.
Paris Conference on AI & Digital Ethics, 2023.
NeurIPS Workshop on Cultures in AI/AI in Culture, 2022.

11. **Panel: Fairness in Machine Learning**, *with Sam Corbett-Davies and Susan J. Brison.*
Neukom Institute for Computational Science, Dartmouth College, 2022.
12. **Algorithmic Bias, Generalist Models, and Clinical Medicine.**
Workshop on the Ethics of Influence, University of Oxford, 2022.
13. **The Ethics of Autonomous Vehicles.**
Philosophy and AI Seminar, Meta, 2022.
Center for Human-Compatible AI, UC Berkeley, 2021.
14. **Enabling Fairness in Healthcare through Machine Learning.**
Stereotyping and Medical AI Colloquium, King's College London, 2021.
15. **Digital Manipulation and Mental Integrity**, *with Christopher Burr.*
Workshop on the Philosophy of Online Manipulation, 2021.
16. **A Dilemma for Reasons Additivity**
Political Theory Workshop, Stanford University, 2021.
17. **On Algorithmic Fairness in Medical Practice**, *with Thomas Grote.*
Workshop on Algorithmic Fairness, University of Copenhagen, 2020.
18. **Decision Support Systems and Clinical Reasoning**, *with Rune Nyrup.*
Leverhulme Centre for the Future of Intelligence, University of Cambridge, 2020.
Philosophy of Medical AI Workshop, University of Tübingen, 2020.
Issues in Explainable AI, Saarland University, 2019.
19. **Proper Functions: Etiology Without Typehood**, *with Niall Paterson.*
9th Philosophy of Biology and Cognitive Sciences Workshop, University of the Basque Country, 2019.
British Society for the Philosophy of Science, University of Durham, 2019.
20. **Object Classification, Autonomous Vehicles, and the Reasonable Belief Standard**
2nd CEPPA Graduate Conference on Moral and Political Philosophy, University of St Andrews, 2019.
Automated Vehicles Symposium, Orlando FL, 2019.
Research Institute for Ethics and Law, Swansea University, 2019.
iCog5: Approaches to Higher Cognitive Function, University of Reading, 2019.
Smart Cities Workshop, University of Bristol, 2019.
21. **Uncertainty About Persons in Autonomous Vehicle Collisions**
Automated Vehicles Symposium, San Francisco, CA, 2018.
22. **Panel: Ethical and Social Implications of Autonomous Vehicles**, *with Nicholas G. Evans, Noah Goodall, Ryan Jenkins, and Katherine Evans.*
Automated Vehicles Symposium, San Francisco, CA, 2018.
23. **Against Leben's Rawlsian Algorithm for Autonomous Vehicles**
Philosophy and Theory of Artificial Intelligence, University of Leeds, 2017.
Artificial Ethics Symposium, University of Southampton, 2017.

24. **Reliability Weighted Belief Revision in Peer Disagreements**, *with Nemo D'Qrill*.
 Graduate Conference on Epistemology of Disagreement, University of Tartu, 2017.
 9th European Congress of Analytic Philosophy, LMU Munich, 2017.
25. **Blame in Autonomous Vehicle Collisions**
 10th Symposium on Computing and Philosophy, Annual Convention of the Society for the Study of Artificial Intelligence and the Simulation of Behaviour, University of Bath, 2017.

Service

Organising Committee , <i>ICLR Workshop on Human-AI Coevolution</i> .	2025
Organising Committee , <i>Paris Conference on AI & Digital Ethics</i> .	2024
Program Committee , AI Alignment Track, <i>AAAI Conference on Artificial Intelligence</i> .	2024
Program Committee , FAccT, <i>Association for Computing Machinery</i> .	2022-2023
Program Committee , Explanatory AI: Between Ethics and Epistemology, <i>TU Delft</i> .	2021-2022
Program Committee , Bias and Discrimination in Algorithmic Decisions, <i>Leibniz University, Hannover</i> .	2021-2022
Program Committee , Philosophy After AI Symposium, <i>AISB</i> .	2019-2020
President , <i>British Postgraduate Philosophical Association</i> .	2018-2019
Executive Committee Member , <i>British Philosophical Association</i> .	2018-2019

Teaching

Ethics, Law, and Politics of Artificial Intelligence , Primary Instructor, <i>Stanford University</i> .	2020-2021
Political Philosophy , Seminar Tutor, <i>University of Bristol</i> .	2019-2020
Research Methods for Social Science , Seminar Tutor and Guest Lecturer, <i>University of Bristol</i> .	2018-2019
Formal Logic , Seminar Tutor, <i>University of Bristol</i> .	2017-2018

Ad-Hoc Teaching

Assessing Consciousness in LLMs , Guest Lecture, <i>Hong Kong University</i> .	2025
Fairness in Machine Learning , Guest Lecture, <i>Princeton University</i> .	2024
Do LLMs have credences? , Guest Lecture, <i>Drake University</i> , with Winnie Street.	2024
Moral Imagination for Responsible Innovation , Guest Lecture, <i>LMU Munich</i> , with Benjamin Lange.	2023
Fairness in Machine Learning , Guest Lecture, <i>Princeton University</i> .	2023
Black Boxes and Explainability , Guest Lecture, <i>Chinese University of Hong Kong</i> .	2021
Ethics of Medical AI , Guest Lecture, <i>Utrecht University</i> .	2020

Awards

PhD Studentship , <i>Arts and Humanities Research Council</i> , c.a. £55,000. Awarded through the South, West and Wales Doctoral Training Partnership.	2017-2020
Andrea Mannu Prize , <i>London School of Economics</i> . Awarded for joint-best performance across all undergraduate philosophy programs (joint-1st/51).	2016

Research Students

Lize Alberts, DPhil Candidate in Computer Science, *University of Oxford*. 2023
Primary supervisor for 3 month Student Researcher engagement at Google Research.

Public Policy

Law Commission of England and Wales and Scottish Law Commission. 2021
Consulted for the 2022 joint report on 'Autonomous Vehicles.' [\[Link\]](#)

Wider Engagement

Fireside Chat: AI Cognition and Consciousness, with *Henry Shevlin*. 2025
Homerton College, University of Cambridge.

Panel: Industry-Academia Collaboration in Responsible AI, with *Murray Shanahan & Winnie Street*. 2024
Leverhulme Centre for the Future of Intelligence, University of Cambridge.

Talk: Thinking about AI Consciousness, with *Winnie Street*. 2024
LSE Philosophy Society, London School of Economics.

Panel: AI, Tech and Humanities Careers, with *Harriet Walker, Dora Szabo & Genevieve Liveley*. 2024
South, West and Wales Doctoral Training Partnership.

Panel: Careers in Philosophy, with *Leticia Garcia Martinez, Hamza King, Rachel Ghaw & Jefferson Courtney*. 2024
Department of Philosophy, Logic, and Scientific Method, London School of Economics.

Panel: Careers in AI Ethics and Policy, with *Risto Uuk & Asbyana-Jasmine Kachra*. 2024
London School of Economics.

Podcast: Could AI Undermine Informed Consent?, with *Rune Nystrup & Reid Blackman*. 2024
Ethical Machines Podcast.

Panel: Spotlight on Ethical AI, with *Mandeep Soor, Ismael Garcia & Timothy Wu*. 2023
London School of Economics.

Talk: Ethics and Large Language Models. 2023
Penningtons Manches Cooper & Ethical Reading.

Talk: Practical Ethics in Silicon Valley. 2022
Uehiro Centre for Practical Ethics, University of Oxford.

Talk: Teaching AI Ethics: Lessons from Stanford and Google. 2022
National Symposium on Developing Socially Responsible STEM Professionals, City, University of London.

Talk: Practical Ethics in Tech. 2021
Let's Phi: Beyond Academic Philosophy.

Talk: AI Ethics and the Long Term Future. 2021
CreAtivity Workshop (for underprivileged high school students).

Talk: Philosophy Careers Outside Academia. 2021
Department of Philosophy, University of Bristol.

Talk: Medical Reasoning in the Age of Artificial Intelligence. 2020
Milton Keynes Artificial Intelligence Expert Forum.

Panel: Artificial Intelligence in Healthcare, with Kourosh Saeb-Parsy, Slawomir Nasuto & Weizi Vicky Li. 2019

Austin Vita & Ethical Reading Public Event on AI in Healthcare.

Talk: The Ethics of Automated Vehicles. 2018

St Edward's School, Oxford.

Peer Review

Referee for *Nature*, *Nature Machine Intelligence*, *Philosophical Studies*, *British Journal for the Philosophy of Science*, *Journal of Applied Philosophy*, *Utilitas*, *Journal of Ethics*, *Synthese*, *Journal of Medical Ethics*, *Ethics and Information Technology*, *Science and Engineering Ethics*, *Philosophy and Technology*, *Res Publica*, *AI and Society*, *IEEE Transactions on Technology and Society*, *IEEE Intelligent Transportation Systems Transactions*, *IEEE Access*, and *Virtual Reality*. I have also refereed for ACM FAccT, NeurIPS, ICLR, and AAAI/ACM Conference on AI, Ethics and Society, alongside the MIT Press and the John Templeton Foundation.