

Geoff Keeling

📍 London, United Kingdom
💡 Ethics of AI/ML, Digital Minds, Machine Psychology,
Bioethics, Moral & Political Philosophy, Cognitive Science

in /in/geoff-keeling
🌐 geoffkeeling.github.io
✉ geoffkeelingbhyat@gmail.com

Employment

Staff Research Scientist , <i>Google, Paradigms of Intelligence Team.</i>	2025–Present
Senior Research Scientist , <i>Google, Paradigms of Intelligence Team.</i>	2023–25
Research Scientist , <i>Google Research.</i>	2023–23
Bioethicist , <i>Google Health.</i>	2021–23
Postdoctoral Fellow , <i>Institute for Human-Centered AI & Center for Ethics in Society, Stanford University.</i>	2020–21
Research Assistant , <i>Leverhulme Centre for the Future of Intelligence, University of Cambridge.</i>	2019–20

Affiliations

Fellow , <i>Institute of Philosophy, School of Advanced Study, University of London.</i>	2024–Present
Associate Fellow , <i>Leverhulme Centre for the Future of Intelligence, University of Cambridge.</i>	2020–Present

Education

PhD Philosophy , <i>University of Bristol</i> , No Corrections.	2017–20
MA Philosophy , <i>University of Bristol</i> , Distinction in all Components (Rank: 1st/24).	2016–17
BSc Philosophy, Logic, and Scientific Method , <i>London School of Economics</i> , First Class (Rank: Joint 1st/51).	2013–16

Book

Keeling, G. & Street, W. (2026). *Emerging questions in AI welfare*. Cambridge University Press.

Papers

- Kim, J., Street, W., Rocca, R., Korngiebel, D., Waytz, A., Evans, J., & **Keeling, G.** (2026). Theory of mind and self-attributions of mentality are dissociable in LLMs. *arXiv:2603.28925v1*.
- Rocca, R., Street, W., **Keeling, G.**, & Evans, J. (2026) Psychological coupling: The necessary science of human-AI interaction. *PsyArXiv:695gp_v1*.
- Marchal, N., Chan, S., Franklin, M., Revel, M., **Keeling, G.**, Fischli, R., Chandra, B., & Gabriel, I. (2026). Architecting trust in artificial epistemic agents. *arXiv:2603.02960*.
- Caviola, L., **Keeling, G.**, Street, W., & Shevlin, H. (2026). Human-AI coexistence. *PhilArchive:CAVHCQ*.
- Keeling, G.** & Street, W. (2026). Chuck, Wilson and the emergence of artificial minds in human-AI conversations. *Journal of Consciousness Studies*. (Accepted Manuscript) [Preprint: *arXiv:2601.13081*.]
- Lange, B., **Keeling, G.**, Pedersen, K., Heringer, C., Rubin, S. B., Zevenbergen, B., & McCroskery, A. (2026). Epistemic trust as a mechanism for ethics integration: Failure modes and design principles from 70 moral imagination workshops. *Journal of Responsible Innovation*. (Accepted Manuscript) [Preprint: *arXiv preprint arXiv:2604.11281v1*.]

7. Street, W., Siy, J.O., **Keeling, G.**, Baranes, A., Barnett, B., McKibben, M., Kanyere, T., Lentz, A., Agüera y Arcas, B., & Dunbar, R. (2026). LLMs achieve adult human performance on higher-order theory of mind tasks. *Frontiers in Human Neuroscience*, 19.
8. Gabriel, I., **Keeling, G.**, Manzini, A., & Evans, J. (2025). We need a new ethics for a world of AI agents. *Nature*, 644, 38-40.
9. Gabriel, I. & **Keeling, G.** (2025). A matter of principle? AI alignment as the fair treatment of claims. *Philosophical Studies*, 182, 1951–1973.
10. **Keeling, G.** & Street, W. (2025). On the attribution of confidence to large language models. *Inquiry*, 1-27.
11. Lange, B., **Keeling, G.**, Manzini, A., & McCroskery, A. (2025). We need accountability in human-AI agent relationships. *npj Artificial Intelligence*, 1(38), 1-3.
12. Grzankowski, A.*, **Keeling, G.***, Shevlin, H.*, & Street, W.* (2025). Deflating deflationism: a critical perspective on debunking arguments against LLM mentality. *arXiv:2506.13403*. (*Equal Contribution)
13. Broestl, N.*, Lange, B.*, Voinea, C.*, **Keeling, G.***, & Lam, R. (2025). Evaluating intra-firm LLM alignment strategies in business contexts. *arXiv:2505.18779v1*. (*Equal Contribution) [Media: [AI Insider](#)]
14. Gabriel, I.*, Manzini, A.*, **Keeling, G.***, Hendricks, L. A., Reiser, V., Iqbal, H., Tomasev, N., Ktena, I., Kenton, Z., Rodriguez, M., El-Sayed, S., Brown, S., Akbulut, C., Trask, A., Hughes, E., Bergman, A. S., Shelby, R., Marchal, N., Griffin, C., Mateos-Garcia, J., Weidinger, L., Street, W., Lange, B., Ingerman, A., Lentz, A., Enger, R., Barakat, A., Krakovna, V., Siy, J. O., Kurth-Nelson, Z., McCroskery, A., Bolina, V., Law, H., Shanahan, M., Alberts, L., Balle, B., de Haas, S., Ibitoye, Y., Dafoe, A., Goldberg, B., Kreier, S., Reese, A., Witherspoon, S., Hawkins, Rauh, M., Wallace, D., Franklin, M., Goldstein, J.A., Lehman, J., Klenk, M., Vallor, S., Biles, C., Ringel Morris, M., King, H., Agüera y Arcas, B., Isaac, W., & Manyika, J. (2024). The ethics of advanced AI assistants. *arXiv:2404.16244*. (*Equal Contribution) [Media: [WIRED](#), [CBS](#), [The Information](#), [Axios](#), [The Verge](#), [MSN](#), [Tech Policy Press](#), [Psychology Today](#), [MarkTechPost](#)]
15. Manzini, A., **Keeling, G.**, Marchal, N., McKee, K., Reiser, V., Haas, J., & Gabriel, I. (2024, June). Should users trust advanced AI assistants? Justified trust as a function of competence and alignment. In: *Proceedings of FAccT '24: ACM Conference on Fairness, Accountability, and Transparency*.
16. Manzini, A., **Keeling, G.**, Alberts, L., Vallor, S., Ringel Morris, M., & Gabriel, I. (2024, October). The code that binds us: Navigating the appropriateness of human-AI assistant relationships. In: *Proceedings of the Seventh AAAI/ACM Conference on AI, Ethics, and Society*. [🏆 Runner Up: Best Paper Award]
17. **Keeling, G.***, Street, W.*, Stachaczyk, M., Zakharova, D., Comşa, I., Sakovych, A., Logothetis, I., Zhang, Z., Agüera y Arcas, B., & Birch, J. (2024). Can LLMs make trade-offs involving stipulated pain and pleasure states? *arXiv:2411.02432v1*. (*Equal Contribution) [Media: [Scientific American](#), [Futurism](#), [Aeon](#), [The Guardian](#)]
18. El-Sayed, S., Akbulut, C., McCroskery, A., **Keeling, G.**, Kenton, Z., Jalan, Z., Marchal, N., Manzini, A., Shevlane, T., Vallor, S., Susser, D., Franklin, M., Bridgers, S., Law, H., Rahtz, M., Shanahan, M., Tessler, M.H., Douillard, A., Everitt, T., & Brown, S. (2024). A mechanism-based approach to mitigating harms from persuasive generative AI. *arXiv:2404.15058*. [Media: [VentureBeat](#), [Futurism](#)]
19. Schaeckermann, M., Spitz, T., Pyles, M., Cole-Lewis, H., Wulczyn, E., Pfhol, S.R., Martin, D., Jaroensri, R., **Keeling, G.**, Liu, Y., Farquhar, S., Xue, Q., Lester, J., Hughes, C., Strachan, P., Tan, F., Bui, P., Mermel, C.H., Peng, L.H., Matias, Y., Corrado, G.S., Webster, D.R., Virmani, S., Semturs, C., Liu, Y., Horn, I., & Chen, P-H.C. (2024). Health equity assessment of machine learning performance (HEAL): a framework and dermatology AI model case study. *Lancet eClinicalMedicine*. [Media: [MarkTechPost](#)]
20. Alberts, L., **Keeling, G.**, & McCroskery, A. (2024). Should agentic conversational AI change how we think about ethics? Characterising an interactional ethics centred on respect. *arXiv:2401.09082v2*.
21. **Keeling, G.**, Lange, B., McCroskery, A., Pedersen, K., Weinberger, D., & Zevenbergen, B. (2024). Moral imagination for engineering teams: The technomoral scenario. *International Review of Information Ethics*, 34(1).

22. Lange, B.*, **Keeling, G.***, McCroskery, A., Zevenbergen, B., Blascovich, S., Pedersen, K., Lentz, A., & Agüera y Arcas, B. (2023). Engaging engineering teams through moral imagination: a bottom-up approach for responsible innovation and ethical culture change in technology companies. *AI and Ethics*, 1-10. (*Equal Contribution)
23. **Keeling, G.** (2023). Algorithmic bias, generalist models, and clinical medicine. *AI and Ethics*, 1-12.
24. **Keeling, G.** & Nyrup R. Explainable machine learning, patient autonomy, and clinical reasoning. In: Véliz, C. (Ed). (2023). *Oxford Handbook of Digital Ethics* (pp. 528–552). Oxford University Press.
25. **Keeling, G.** (2023). A dilemma for reasons additivity. *Economics & Philosophy*, 39(1), 20-42.
26. Grote, T. & **Keeling, G.** (2022). Enabling fairness in healthcare through machine learning. *Ethics and Information Technology*, 24(3), 39.
27. **Keeling, G.** Automated vehicles and the ethics of classification. In: Jenkins, R., Cerný, D., & Hříbek, T. (Eds). (2022). *Autonomous Vehicle Ethics: The Trolley Problem and Beyond* (pp. 41-57). Oxford University Press.
28. Grote, T. & **Keeling, G.** (2022). On algorithmic fairness in medical practice. *Cambridge Quarterly of Healthcare Ethics*, 31(1), 83-94.
29. **Keeling, G.** & Burr, C. Digital manipulation and mental integrity. In: Jongepier, F., & Klenk, M. (2022). *The Philosophy of Online Manipulation* (pp. 253-271). Routledge.
30. **Keeling, G.** & Paterson, N. (2022). Proper functions: etiology without typehood. *Biology & Philosophy*, 37(3), 19.
31. Bommasani, R., ... , **Keeling, G.**, ... , & Liang, P. (2021). On the opportunities and risks of foundation models. *arXiv:2108.07258*. [Media: [Forbes](#), [The Economist](#), [VentureBeat](#), [WIRED](#)]
32. **Keeling, G.** (2020). Why trolley problems matter for the ethics of automated vehicles. *Science and Engineering Ethics*, 26(1), 293-307.
33. **Keeling, G.**, Evans, K., Thornton, S. M., Mecacci, G., & Santoni de Sio, F. Four perspectives on what matters for the ethics of automated vehicles. In: Meyer, G. & Beiker, S. (Eds). (2019). *Road Vehicle Automation 6* (pp. 49-60). Springer International Publishing.
34. **Keeling, G.** (2018). Legal necessity, Pareto efficiency & justified killing in autonomous vehicle collisions. *Ethical Theory and Moral Practice*, 21(2), 413-427.
35. **Keeling, G.** Against Leben's Rawlsian collision algorithm for autonomous vehicles. In: Müller, V. (Ed). (2018). *Philosophy and Theory of Artificial Intelligence 2017* (pp. 259-272). Springer International Publishing.
36. **Keeling, G.** (2018). The sensitivity argument against child euthanasia. *Journal of Medical Ethics*, 44(2), 143-144.
37. **Keeling, G.** (2018). Autonomy, nudging and post-truth politics. *Journal of Medical Ethics*, 44(10), 721-722.
38. **Keeling, G.** (2017). Commentary: Using virtual reality to assess ethical decisions in road traffic scenarios: applicability of value-of-life-based models and influences of time pressure. *Frontiers in Behavioral Neuroscience*, 11, 247-48.

Talks

Chuck, Wilson and the emergence of artificial minds in human-AI conversations, with *Winnie Street*.

AI Minds Symposium, Institute of Philosophy, University of London, 2026.

Panel: Welfare and Moral Status under Uncertainty, with *Andy Stawasz and Winnie Street*.

Center for Mind, Ethics and Policy Summit, New York University, 2026.

Emerging Questions in AI Welfare, with *Winnie Street*.

Book Symposium, Department of Philosophy, Hong Kong University, 2025.

[with replies from Herman Cappelen, Simon Goldstein, Hua Shen, Boris Babic and Nate Sharadin]

Summer School on Consciousness and Metacognition, Institute of Cognitive Neuroscience, UCL, 2025.

Center for Mind, Ethics and Policy Series on AI Welfare, New York University, 2025.

Eleos Conference on AI Consciousness and Welfare, Berkeley, 2025.

Social Contract Research Network, Monash University, 2026.

Book Symposium, Institute of Philosophy, University of London, 2026.

[with replies from Henry Shevlin, Lucia Meloni and Patrick Butlin.]

Panel: AI and Moral Patienthood, *with Henry Shevlin and Winnie Street*.

AI, Animals and Digital Minds Conference, London, 2025.

Can LLMs Make Trade-Offs Involving Stipulated Pain and Pleasure States?, *with Winnie Street*.

Association for the Scientific Study of Consciousness, Heraklion, 2025 (Poster).

AI Welfare Workshop, Eleos AI, Berkeley, 2025.

Institute of Philosophy, University of London, 2024.

The Ethics of Advanced AI Assistants.

AI Ethics, Risks and Safety Conference, Bristol, 2025.

Being Human in the Age of AI, Institute of Philosophy, University of London, 2024.

Ethics of Generative AI & Conversational Agents, LMU Munich, 2024.

Autonomy in an AI World, Imperial College London, 2024.

What About the Moral Claims of AI Agents?

Symposium: Themes from Iason Gabriel, Department of Philosophy, Hong Kong University, 2025.

Do Large Language Models Have Credences?, *with Winnie Street*.

AI and Ethics Workshop, Princeton University, 2024.

Institute for Ethics in Technology, Hamburg University of Technology, 2024.

Educational Impact of Advanced AI Assistants.

AI in Education Workshop, University of Bristol, 2024.

Engaging Google Teams Through Moral Imagination, *with Amanda McCroskey, Ben Zevenbergen, et al.*

Paris Conference on AI & Digital Ethics, 2023.

NeurIPS Workshop on Cultures in AI/AI in Culture, 2022.

Panel: Fairness in Machine Learning, *with Sam Corbett-Davies and Susan J. Brison*.

Neukom Institute for Computational Science, Dartmouth College, 2022.

Algorithmic Bias, Generalist Models, and Clinical Medicine.

Workshop on the Ethics of Influence, University of Oxford, 2022.

Synthetic Characters, Artificial Companionship and Loneliness.

Responsibility and Autonomy in AI, University of Oxford, 2022.

The Ethics of Autonomous Vehicles.

Philosophy and AI Seminar, Meta, 2022.

Center for Human-Compatible AI, UC Berkeley, 2021.

Enabling Fairness in Healthcare through Machine Learning, *with Thomas Grote*.

Stereotyping and Medical AI Colloquium, King's College London, 2021.

Workshop on Algorithmic Fairness, University of Copenhagen, 2020.

Decision Support Systems and Clinical Reasoning, *with Rune Nyrup*.

Leverhulme Centre for the Future of Intelligence, University of Cambridge, 2020.

Philosophy of Medical AI Workshop, University of Tübingen, 2020.

Issues in Explainable AI, Saarland University, 2019.

Digital Manipulation and Mental Integrity, with *Christopher Burr*.

Workshop on the Philosophy of Online Manipulation, 2021.

A Dilemma for Reasons Additivity.

Political Theory Workshop, Stanford University, 2021.

Proper Functions: Etiology Without Typehood, with *Niall Paterson*.

9th Philosophy of Biology and Cognitive Sciences Workshop, University of the Basque Country, 2019.

British Society for the Philosophy of Science, University of Durham, 2019.

Autonomous Vehicles, Object Classification and the Reasonable Belief Standard.

2nd CEPPA Graduate Conference on Moral and Political Philosophy, University of St Andrews, 2019.

Automated Vehicles Symposium, Orlando FL, 2019.

Research Institute for Ethics and Law, Swansea University, 2019.

iCog5: Approaches to Higher Cognitive Function, University of Reading, 2019.

Smart Cities Workshop, University of Bristol, 2019.

Navigating Uncertainty in Autonomous Vehicle Collisions

Automated Vehicles Symposium, San Francisco, CA, 2018.

Panel: Ethics of Autonomous Vehicles, with *Nicholas G. Evans, Noah Goodall, Ryan Jenkins, and Katherine Evans*.

Automated Vehicles Symposium, San Francisco, CA, 2018.

Building Machines that Learn and Think About Morality, with *Christopher Burr*.

Society for the Study of Artificial Intelligence and the Simulation of Behaviour, University of Liverpool, 2018.

Against Leben's Rawlsian Algorithm for Autonomous Vehicles.

Philosophy and Theory of Artificial Intelligence, University of Leeds, 2017.

Reliability Weighted Belief Revision in Peer Disagreements.

Graduate Conference on Epistemology of Disagreement, University of Tartu, 2017.

9th European Congress of Analytic Philosophy, LMU Munich, 2017.

Blame in Autonomous Vehicle Collisions.

Society for the Study of Artificial Intelligence and the Simulation of Behaviour, University of Bath, 2017.

Service

Organising Committee, *CogSci Workshop on Hybrid Cognition in the Era of Human-AI Collaboration*. 2026

Organising Committee, *ICLR Workshop on Human-AI Coevolution*. 2025

Organising Committee, *Paris Conference on AI & Digital Ethics*. 2024

Program Committee, FAccT, *Association for Computing Machinery*. 2022-2023

Program Committee, Explanatory AI: Between Ethics and Epistemology, *TU Delft*. 2021-2022

Program Committee, Bias and Discrimination in Algorithmic Decisions, *Leibniz University, Hannover*. 2021-2022

Program Committee, Philosophy After AI Symposium, *AISB*. 2019-2020

President, *British Postgraduate Philosophical Association*. 2018-2019

Executive Committee Member, *British Philosophical Association*. 2018-2019

Teaching

Ethics, Law, and Politics of Artificial Intelligence, Primary Instructor, *Stanford University*. 2020-2021

Political Philosophy, Seminar Tutor, *University of Bristol*. 2019-2020

Knowledge and Reality , Exam Marker, <i>University of Bristol</i> .	2019-2020
Research Methods for Social Science , Seminar Tutor and Guest Lecturer, <i>University of Bristol</i> .	2018-2019
Formal Logic , Seminar Tutor, <i>University of Bristol</i> .	2017-2018

Ad-Hoc Teaching

AI and Moral Status , Lecture, <i>Venice Philosophy of AI Summer School</i> , with Winnie Street.	2026
Do LLMs have Credences? , Guest Lecture, <i>Hong Kong University</i> , with Winnie Street.	2025
Assessing Consciousness in LLMs , Guest Lecture, <i>Hong Kong University</i> .	2025
Fairness in Machine Learning , Guest Lecture, <i>Princeton University</i> .	2024
Do LLMs have Credences? , Guest Lecture, <i>Drake University</i> , with Winnie Street.	2024
Moral Imagination for Responsible Innovation , Guest Lecture, <i>LMU Munich</i> , with Benjamin Lange.	2023
Fairness in Machine Learning , Guest Lecture, <i>Princeton University</i> .	2023
Black Boxes and Explainability , Guest Lecture, <i>Chinese University of Hong Kong</i> .	2021
Ethics of Medical AI , Guest Lecture, <i>Utrecht University</i> .	2020

Grants and Awards

PhD Studentship , <i>Arts and Humanities Research Council</i> , c.a. £55,000. Awarded through the South, West and Wales Doctoral Training Partnership.	2017-2020
Andrea Mannu Prize , <i>London School of Economics</i> . Awarded for joint-best performance across all undergraduate philosophy programs (joint-1st/51).	2016

Research Students and Mentees

Zandi Eberstadt , DPhil Candidate in Computer Science, <i>University of Oxford</i> . Co-supervisor for Student Researcher placement at Google.	2026-Pres
Zheng (Zack) Hui , PhD Candidate in Computational Linguistics, <i>University of Cambridge</i> . Co-supervisor for Research Intern placement at Google.	2026-Pres
Mariana Amendoeira Duarte , PhD Candidate in Neuroscience, <i>Champalimaud Foundation, Lisbon</i> . Research Mentor for Neuromatch AI Sentience Scholars program.	2026-Pres
Valen Tagliabue , MS Candidate in Artificial Intelligence, <i>University of Milan</i> . Research Mentor for Cambridge Digital Minds Fellowship.	2026-Pres
Junsol Kim , PhD Candidate in Computational Sociology, <i>University of Chicago</i> . Co-supervisor for Student Researcher placement at Google.	2025-Pres
Lize Alberts , DPhil Candidate in Computer Science, <i>University of Oxford</i> . Supervisor for Student Researcher placement at Google.	2023-2023

Public Policy

Research Advisory Board Member , ‘The Delegation Conversation,’ <i>Ada Lovelace Institute</i> .	2026
Consultant , ‘The Dilemmas of Delegation,’ <i>Ada Lovelace Institute</i> .	2025
Consultant , ‘Automated Vehicles,’ <i>Law Commission of England and Wales & Scottish Law Commission</i> .	2021

Wider Engagement

Experience Day: Philosophy, Ethics and AI , with Rev. Ian Jones.	2026
School trip to Google for Sixth Form students at Wymondham College, Norfolk.	
Fireside Chat: AI Consciousness, Welfare and Public Policy , with Jonathan Birch.	2026
Seminar for MSc Philosophy and Public Policy Students, London School of Economics.	
Panel: Philosophical Questions about AI , with Murray Shanahan, Ben Henke & Winnie Street.	2026
British Philosophical Association #PhilosophyMatters Event, University of London.	
Talk: Why Philosophy Matters in AI Development , with Winnie Street.	2026
Institute of Philosophy 20th Anniversary Celebration, School of Advanced Study, University of London.	
Careers Panel , with Anil Seth, Clara Colombatto, Matthias Michel, Tobias Schlicht and Winnie Street.	2025
28th Meeting of the Association for the Scientific Study of Consciousness, Herakleion, Crete.	
Fireside Chat: AI Cognition and Consciousness , with Henry Shevlin.	2025
Seminar for MSt AI Ethics & Society Students, Homerton College, University of Cambridge.	
Panel: Industry-Academia Collaboration in Responsible AI , with Murray Shanahan & Winnie Street.	2024
Leverhulme Centre for the Future of Intelligence, University of Cambridge.	
Talk: Thinking about AI Consciousness , with Winnie Street.	2024
LSE Philosophy Society, London School of Economics.	
Panel: AI, Tech and Humanities Careers , with Harriet Walker, Dora Szabo & Genevieve Liveley.	2024
South, West and Wales Doctoral Training Partnership.	
Panel: Careers in Philosophy , with Leticia García Martínez, Hamza King, Rachel Ghaw & Jefferson Courtney.	2024
Department of Philosophy, Logic, and Scientific Method, London School of Economics.	
Panel: Careers in AI Ethics and Policy , with Risto Uk & Ashyana-Jasmine Kachra.	2024
London School of Economics.	
Podcast: Could AI Undermine Informed Consent? , with Rune Nyrup & Reid Blackman.	2024
Ethical Machines Podcast.	
Panel: Spotlight on Ethical AI , with Mandeep Soor, Ismael Garcia & Timothy Wu.	2023
London School of Economics.	
Talk: Ethics and Large Language Models.	2023
Penningtons Manches Cooper & Ethical Reading.	
Talk: Practical Ethics in Silicon Valley.	2022
Uehiro Centre for Practical Ethics, University of Oxford.	
Talk: Teaching AI Ethics: Lessons from Stanford and Google.	2022
National Symposium on Developing Socially Responsible STEM Professionals, City, University of London.	
Talk: Practical Ethics in Tech.	2021
Let's Phi: Beyond Academic Philosophy.	
Talk: AI Ethics and the Long Term Future.	2021
CreAtivity Workshop (for underprivileged high school students).	
Talk: Philosophy Careers Outside Academia.	2021
Department of Philosophy, University of Bristol.	
Talk: Trolley Problems and Automated Vehicles.	2021
Ethical Reading.	
Talk: Medical Reasoning in the Age of Artificial Intelligence.	2020
Milton Keynes Artificial Intelligence Expert Forum.	

- Panel: Artificial Intelligence in Healthcare**, with Kourosh Saeb-Parsy, Slawomir Nasuto & Weizi Vicky Li. 2019
Austin Vita & Ethical Reading Public Event on AI in Healthcare.
- Talk: The Ethics of Automated Vehicles.** 2018
St Edward's School, Oxford.

Professional Memberships

- Fellow**, Royal Society of Arts. 2025—Pres.

Peer Review

Referee for *Nature*, *Nature Machine Intelligence*, *Scientific Reports*, *Philosophical Studies*, *British Journal for the Philosophy of Science*, *Philosophy of Science*, *Journal of Applied Philosophy*, *Minds and Machines*, *Utilitas*, *Journal of Ethics*, *Synthese*, *Journal of Consciousness Studies*, *Journal of Medical Ethics*, *Ethics and Information Technology*, *Science and Engineering Ethics*, *Philosophy and Technology*, *Res Publica*, *AI and Society*, *AI and Ethics*, *Philosophy and the Mind Sciences*, *IEEE Transactions on Technology and Society*, *IEEE Intelligent Transportation Systems Transactions*, *IEEE Access*, and *Virtual Reality*. I have also refereed for FAccT, NeurIPS, ICLR, AIES, MIT Press, the European Research Council and the John Templeton Foundation.